


What is Apache Spark?

- Open-source distributed computing engine
- Designed for large-scale data processing
- Written in Scala, supports Java, Python (PySpark), and R
- Built around Resilient Distributed Datasets (RDDs)
- Modules: Spark SQL, MLlib, GraphX, Structured Streaming



Features of Apache Spark

- **Speed**
Optimized execution with in-memory computing
- **Ease of Use**
High-level APIs for Java, Scala, Python, and R
- **Libraries**
Built-in modules for SQL, machine learning, streaming and graph processing
- **Scalability**
Capable of scaling out across large clusters